

Harness the power of data in motion with Python and Kafka


Dave Klein

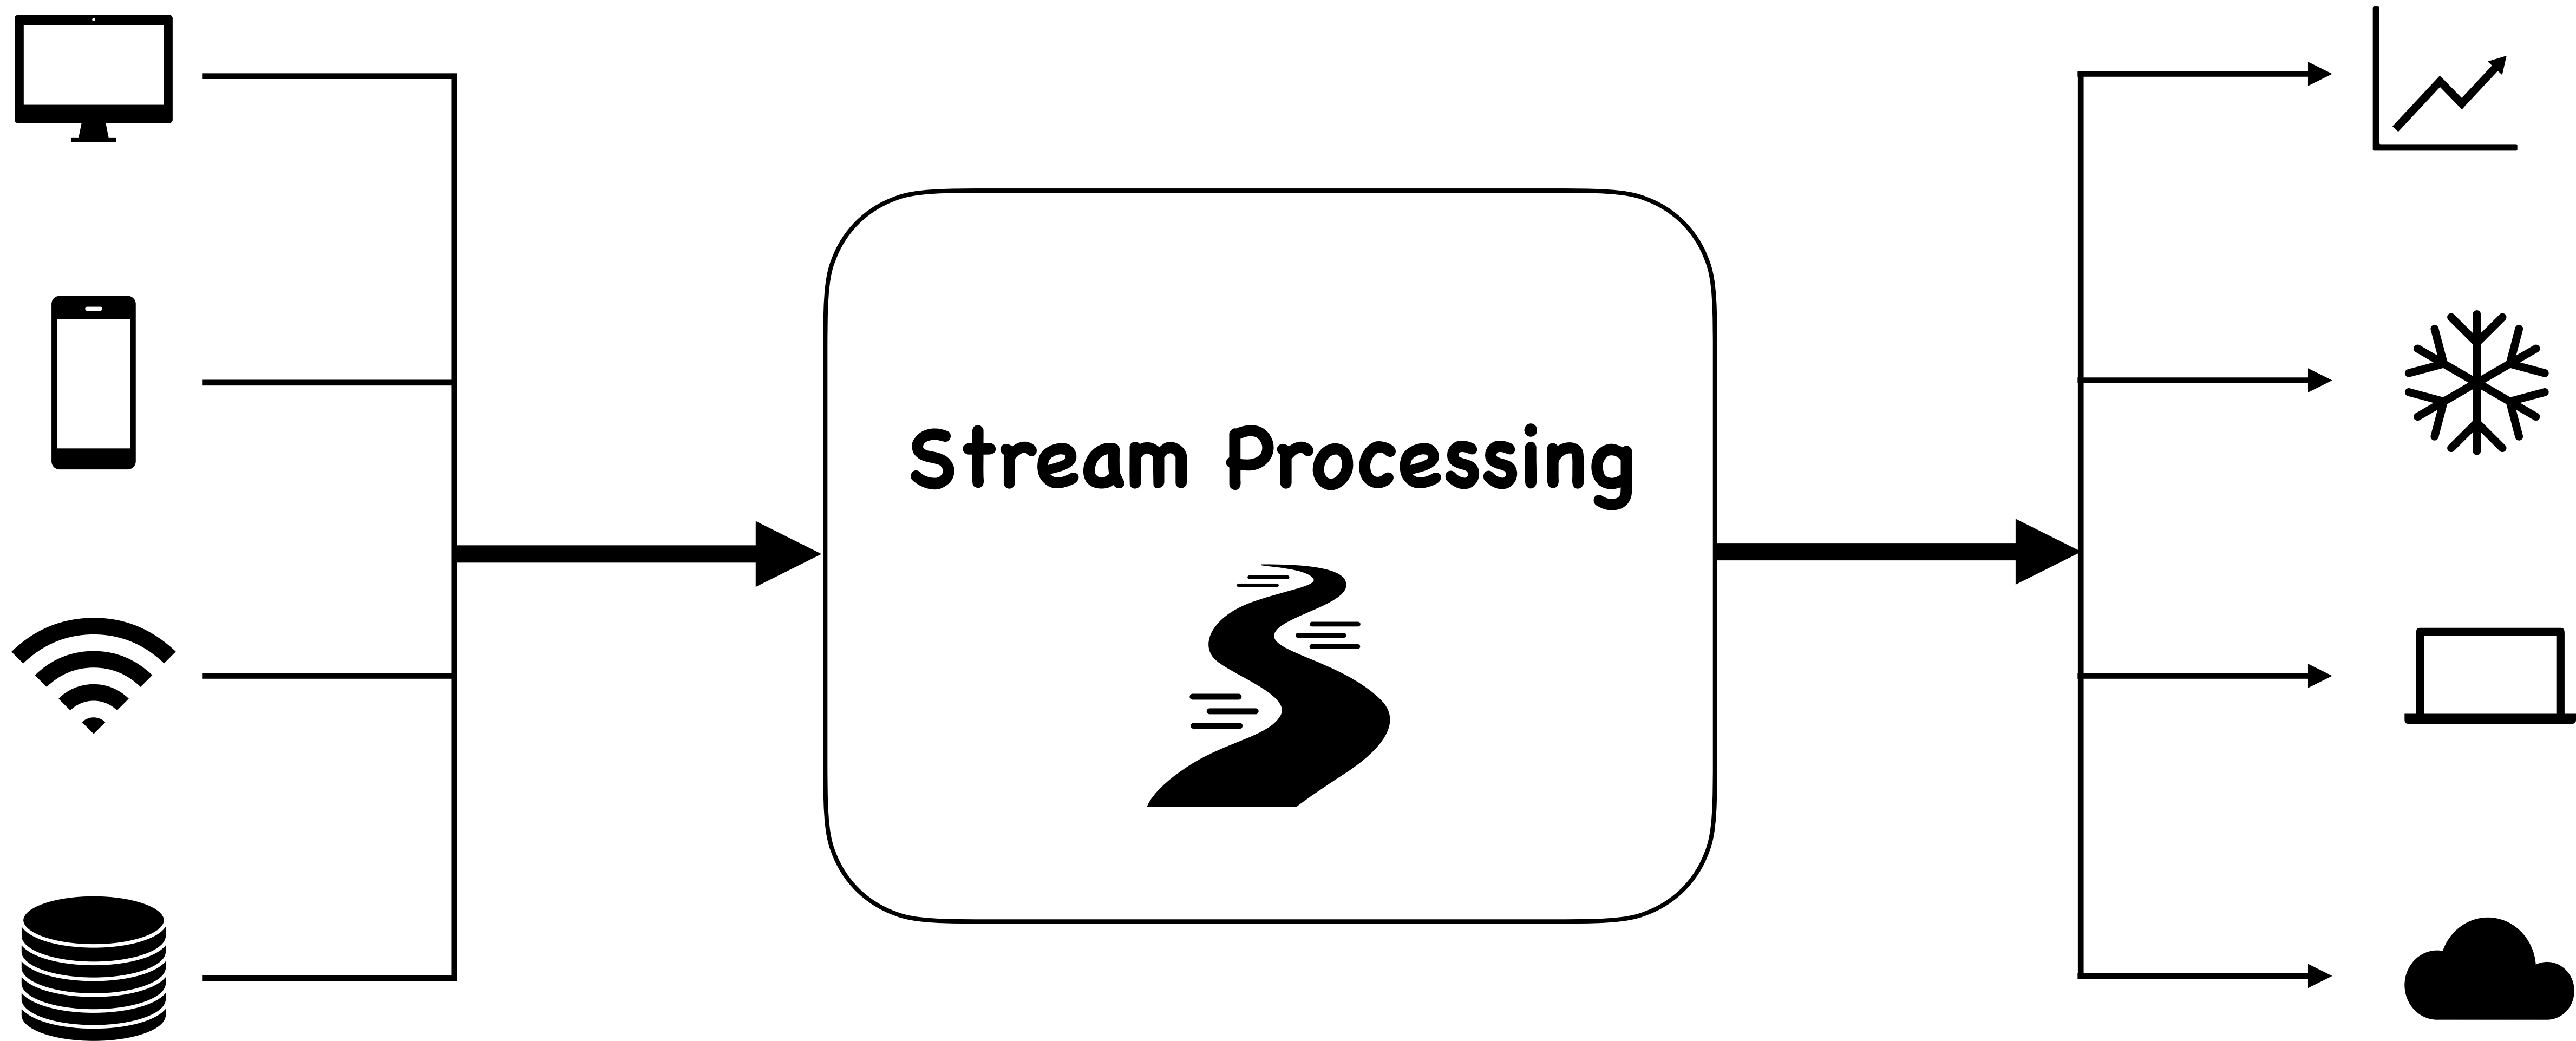
daveklein@usa.net

twitter.com/daveklein

linkedin.com/in/daveklein19

What is Data in Motion?





What Is It Good For?

Customer 360

Fraud Detection

Recommendation Engines


Customer Loyalty Programs

Factory Automation

Fleet Management

Inventory Management


Real-time Payments



Kafka


The Diminutive Guide

Topic (log)




Topic Partitions


P1



P2




P3




Brokers and Replication


Broker 1



Broker 2

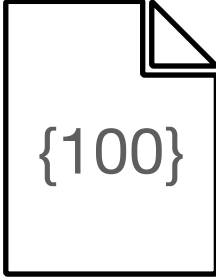
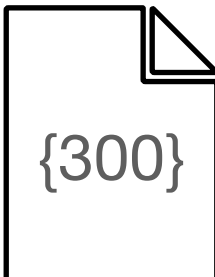
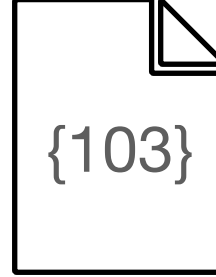




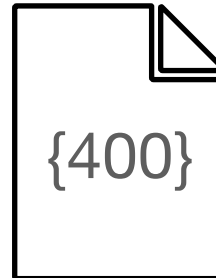
Broker 3


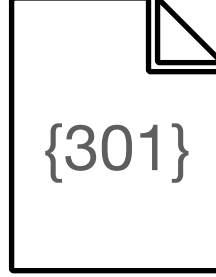




```
"key": "a1",  
"value": {  
  "eventType": "added-to-cart",  
  "title": "Kafka Streams in Action",  
  "author": "Bill Bejeck",  
  "price": 44.99  
}
```

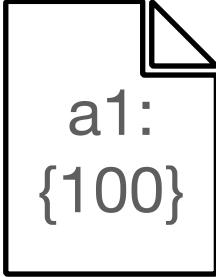
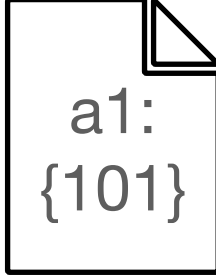

Topic



 {100}	 {300}	 {103}								
1	2	3	4	5	6	7	8	9	10	11

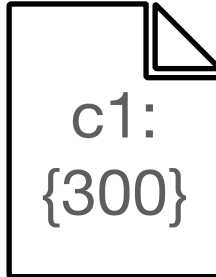
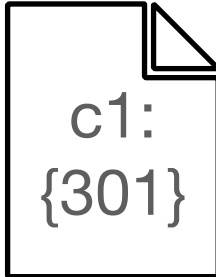
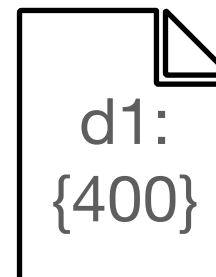

 {200}	 {201}	 {400}								
1	2	3	4	5	6	7	8	9	10	11

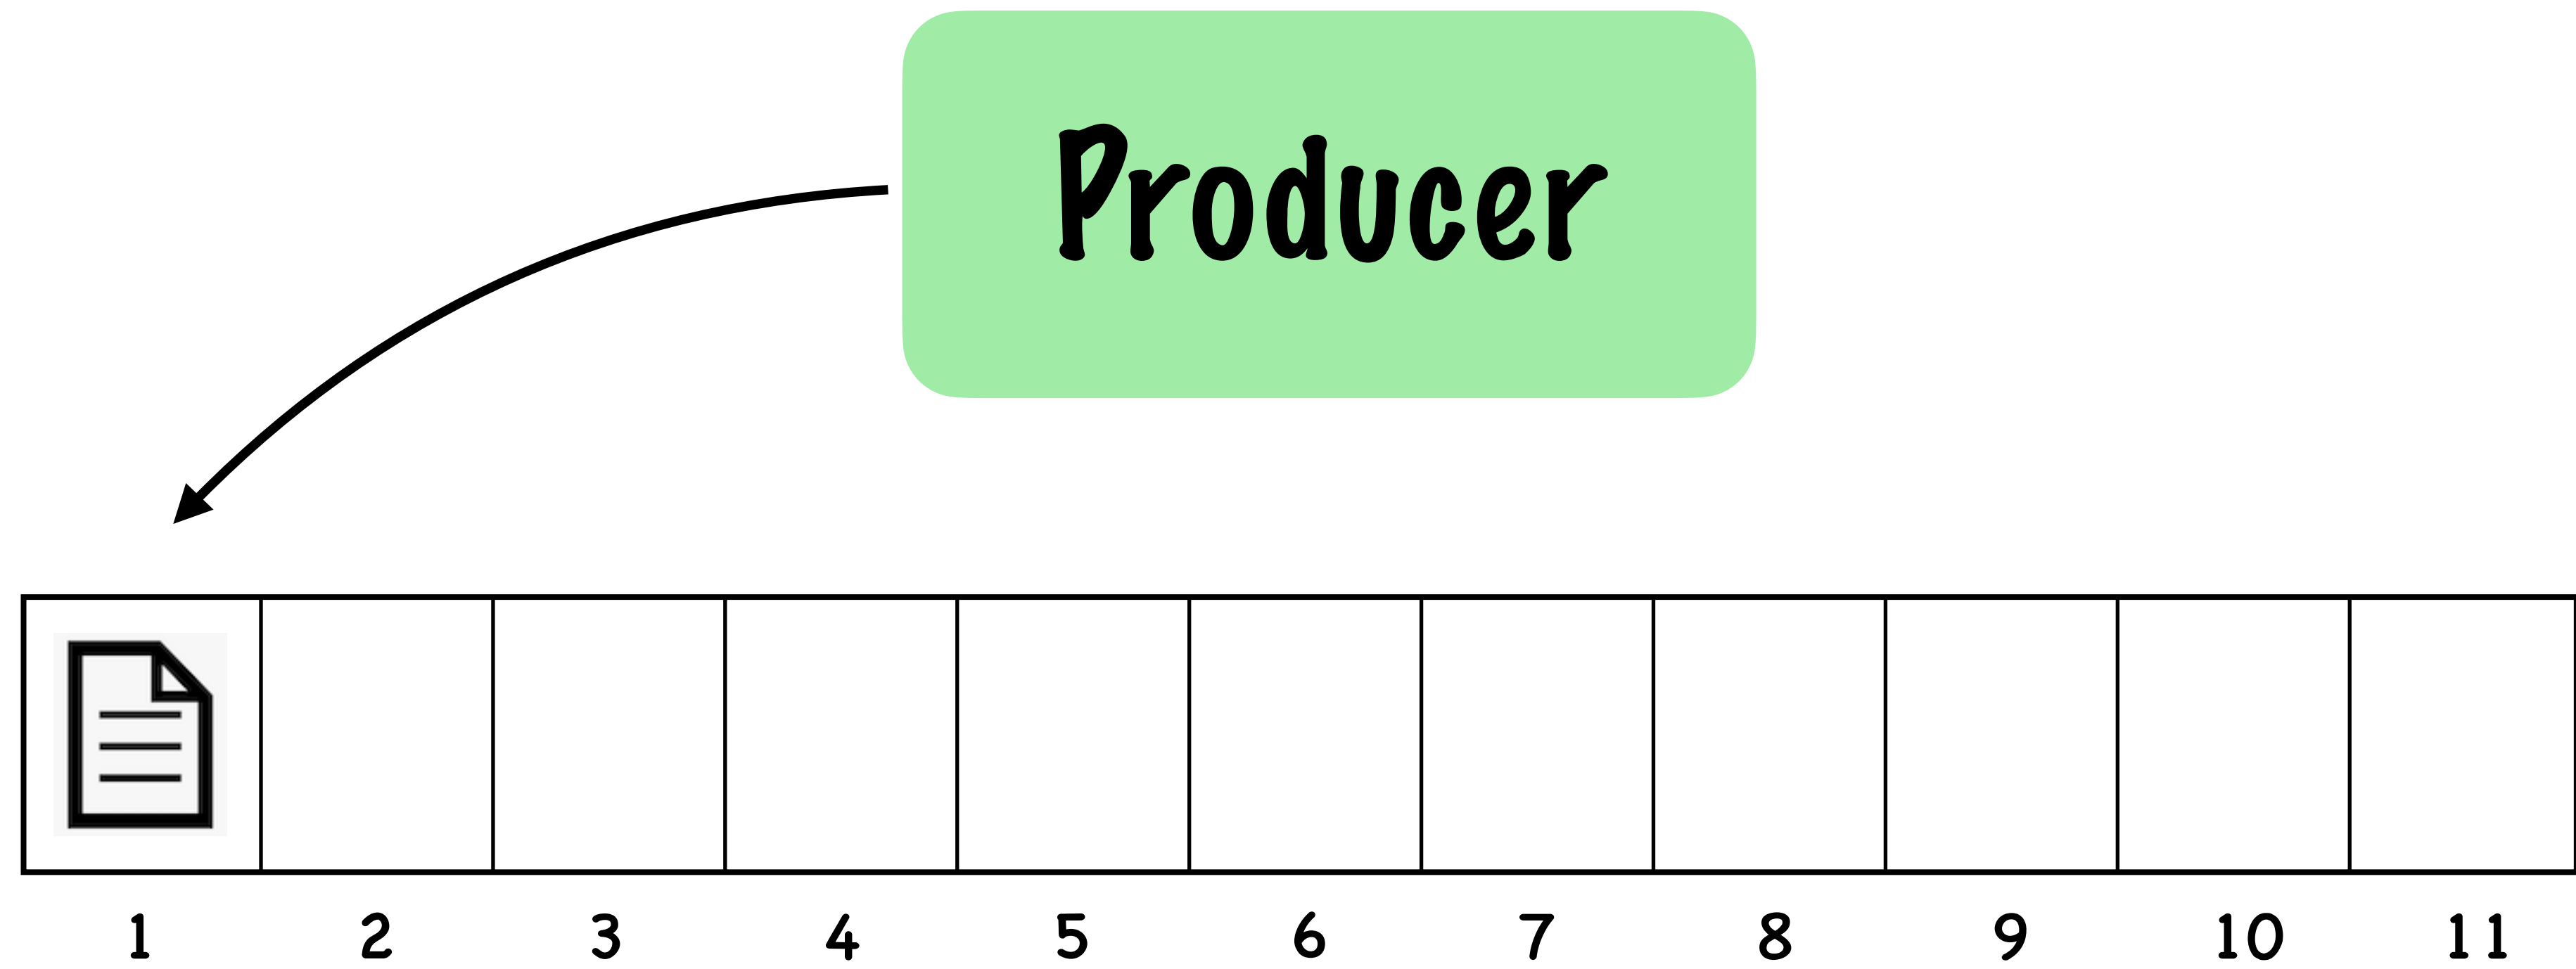
 {101}	 {301}	 {401}								
1	2	3	4	5	6	7	8	9	10	11

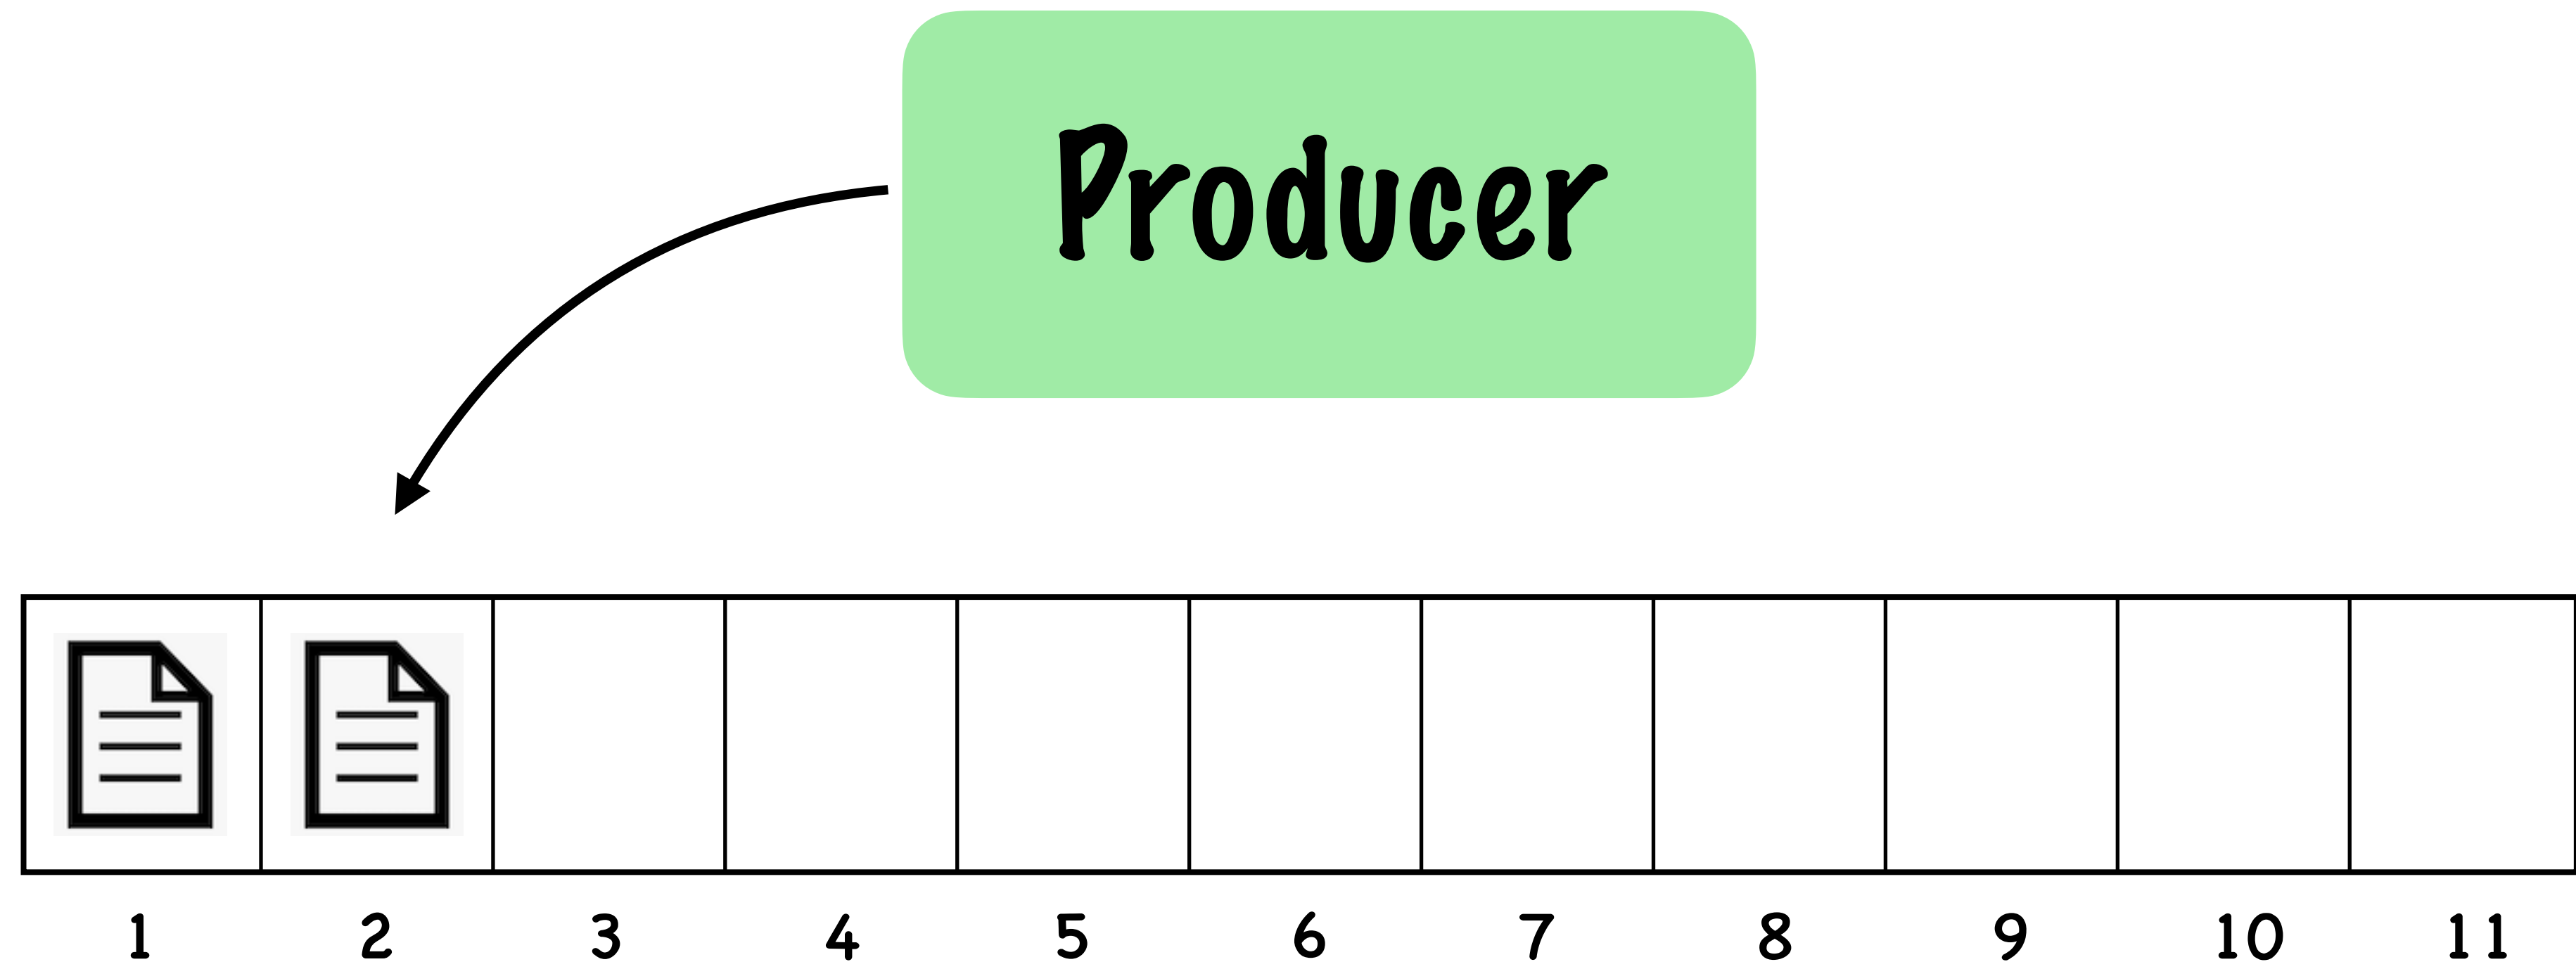
Topic

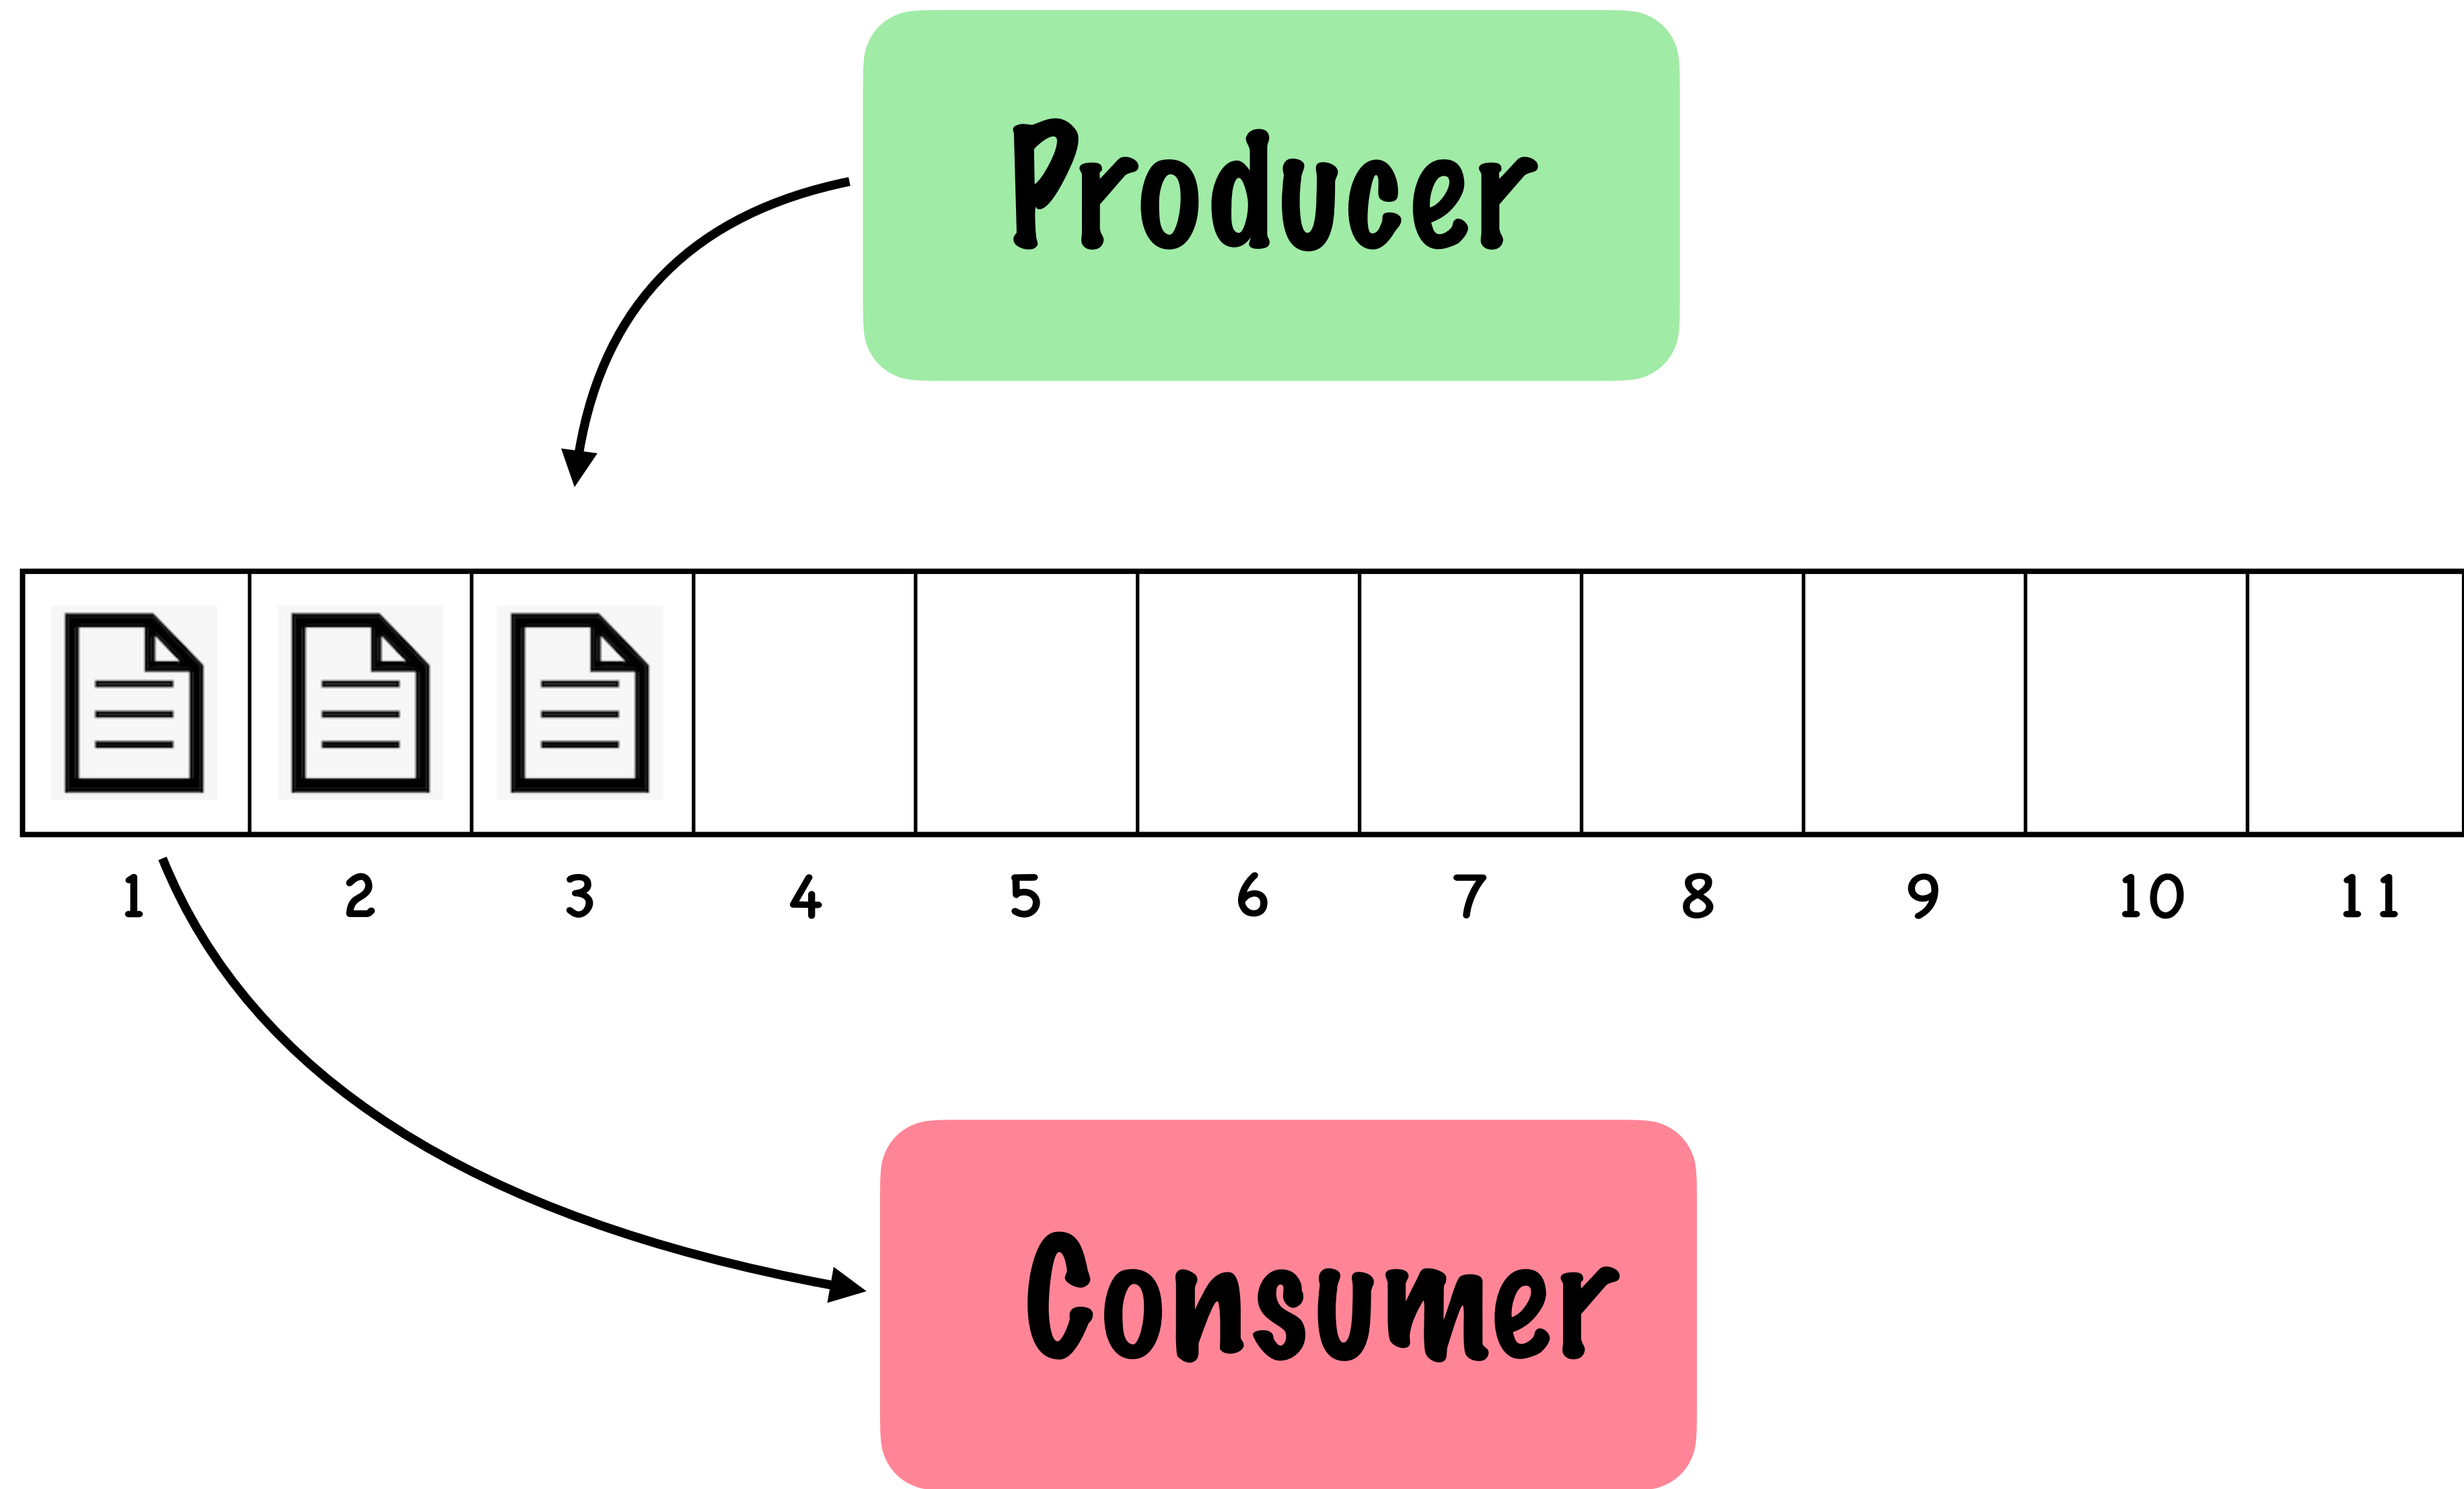
 a1: {100}	 a1: {101}	 a1: {103}								
1	2	3	4	5	6	7	8	9	10	11

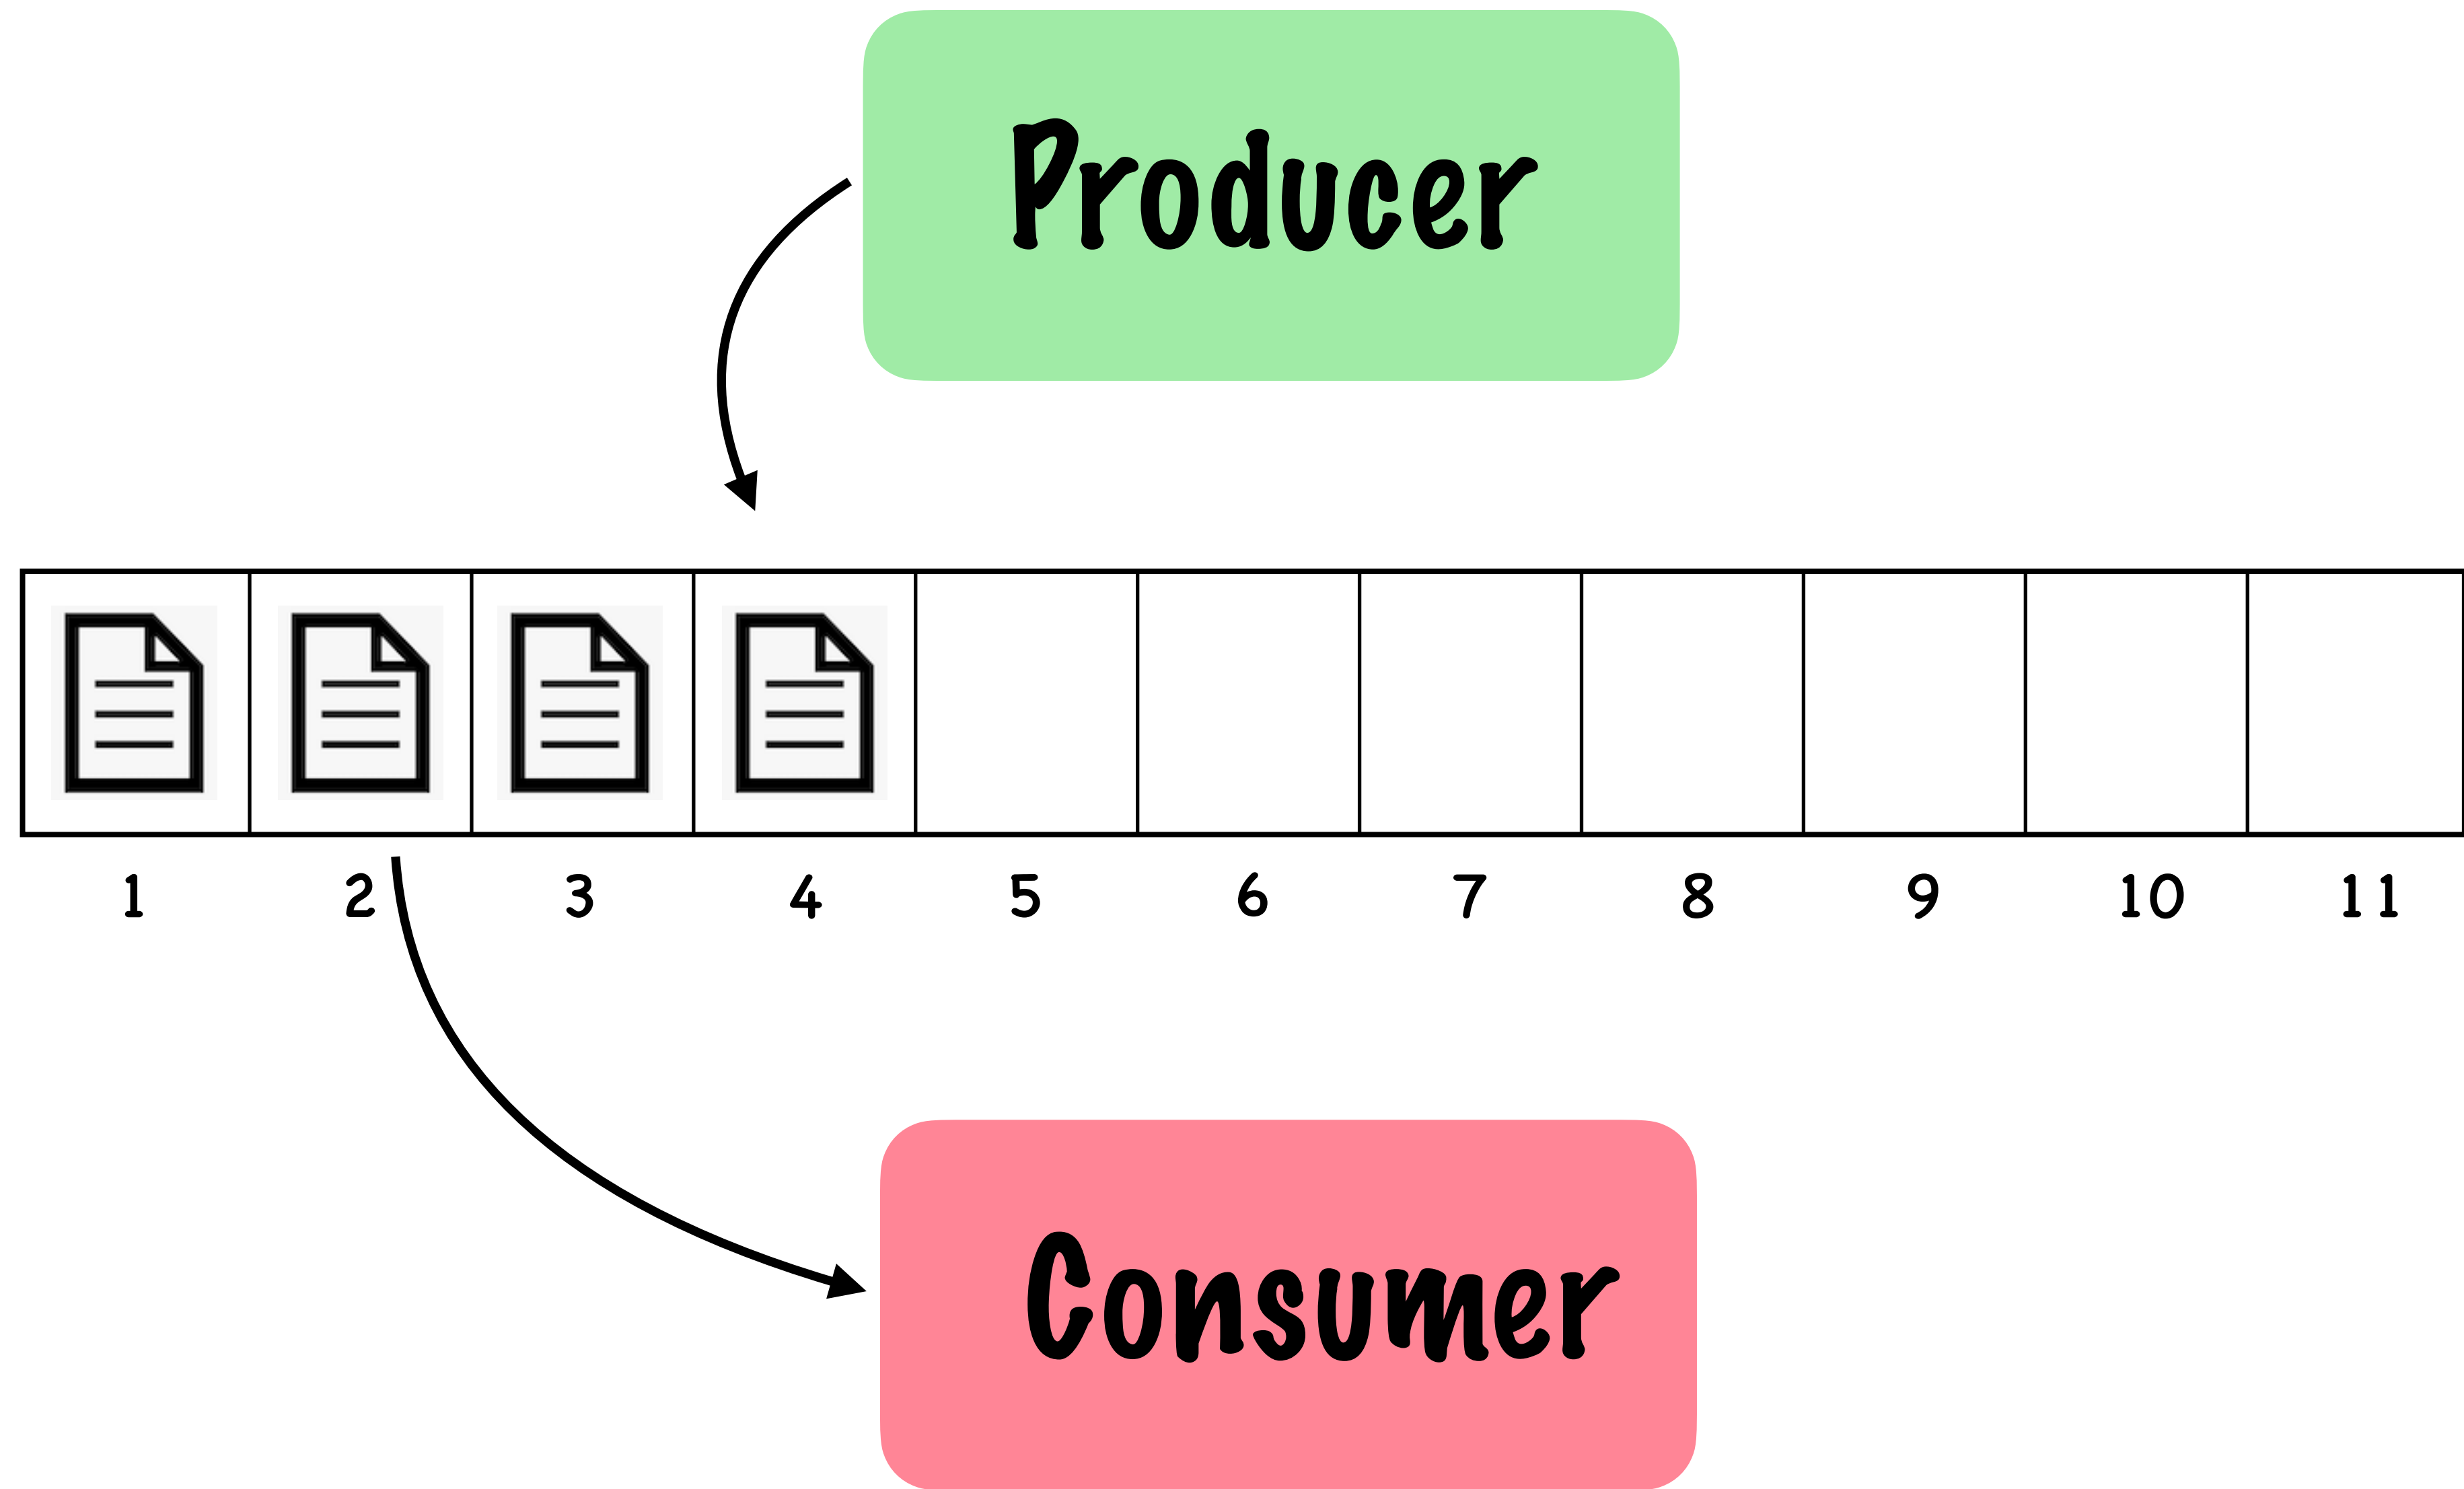
 b1: {200}	 b1: {201}									
1	2	3	4	5	6	7	8	9	10	11

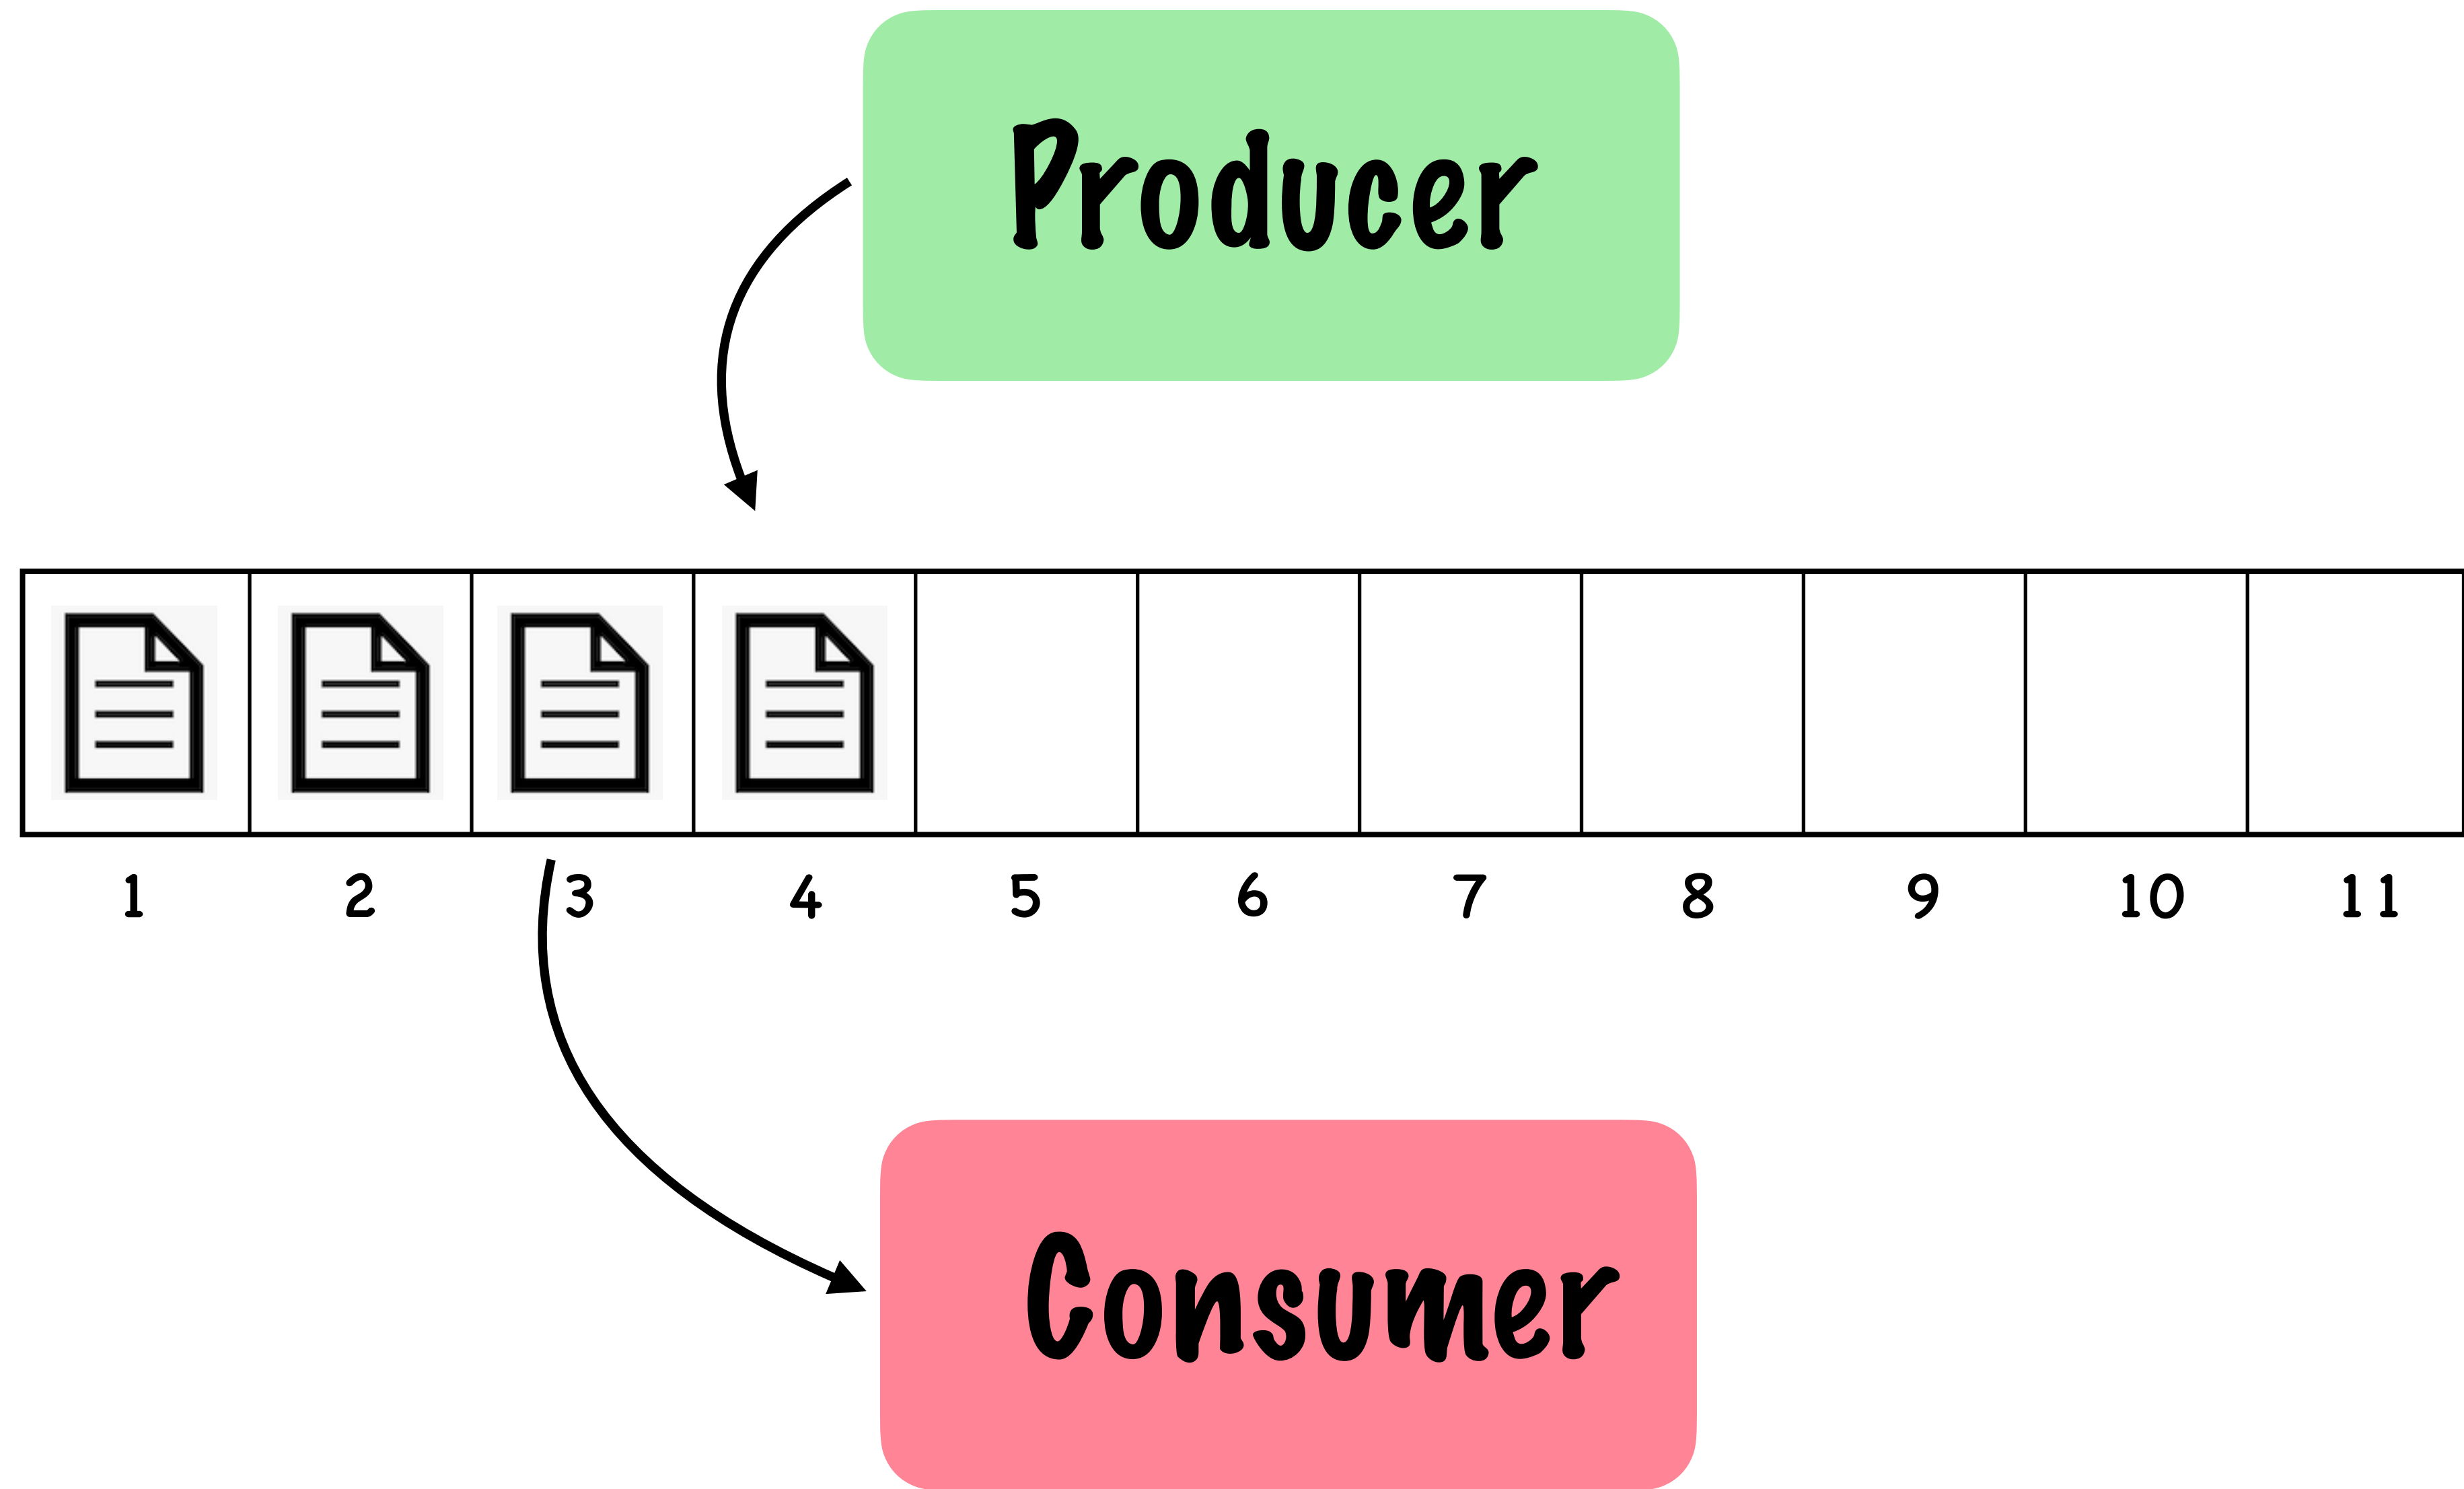
 c1: {300}	 c1: {301}	 d1: {400}	 d1: {401}							
1	2	3	4	5	6	7	8	9	10	11

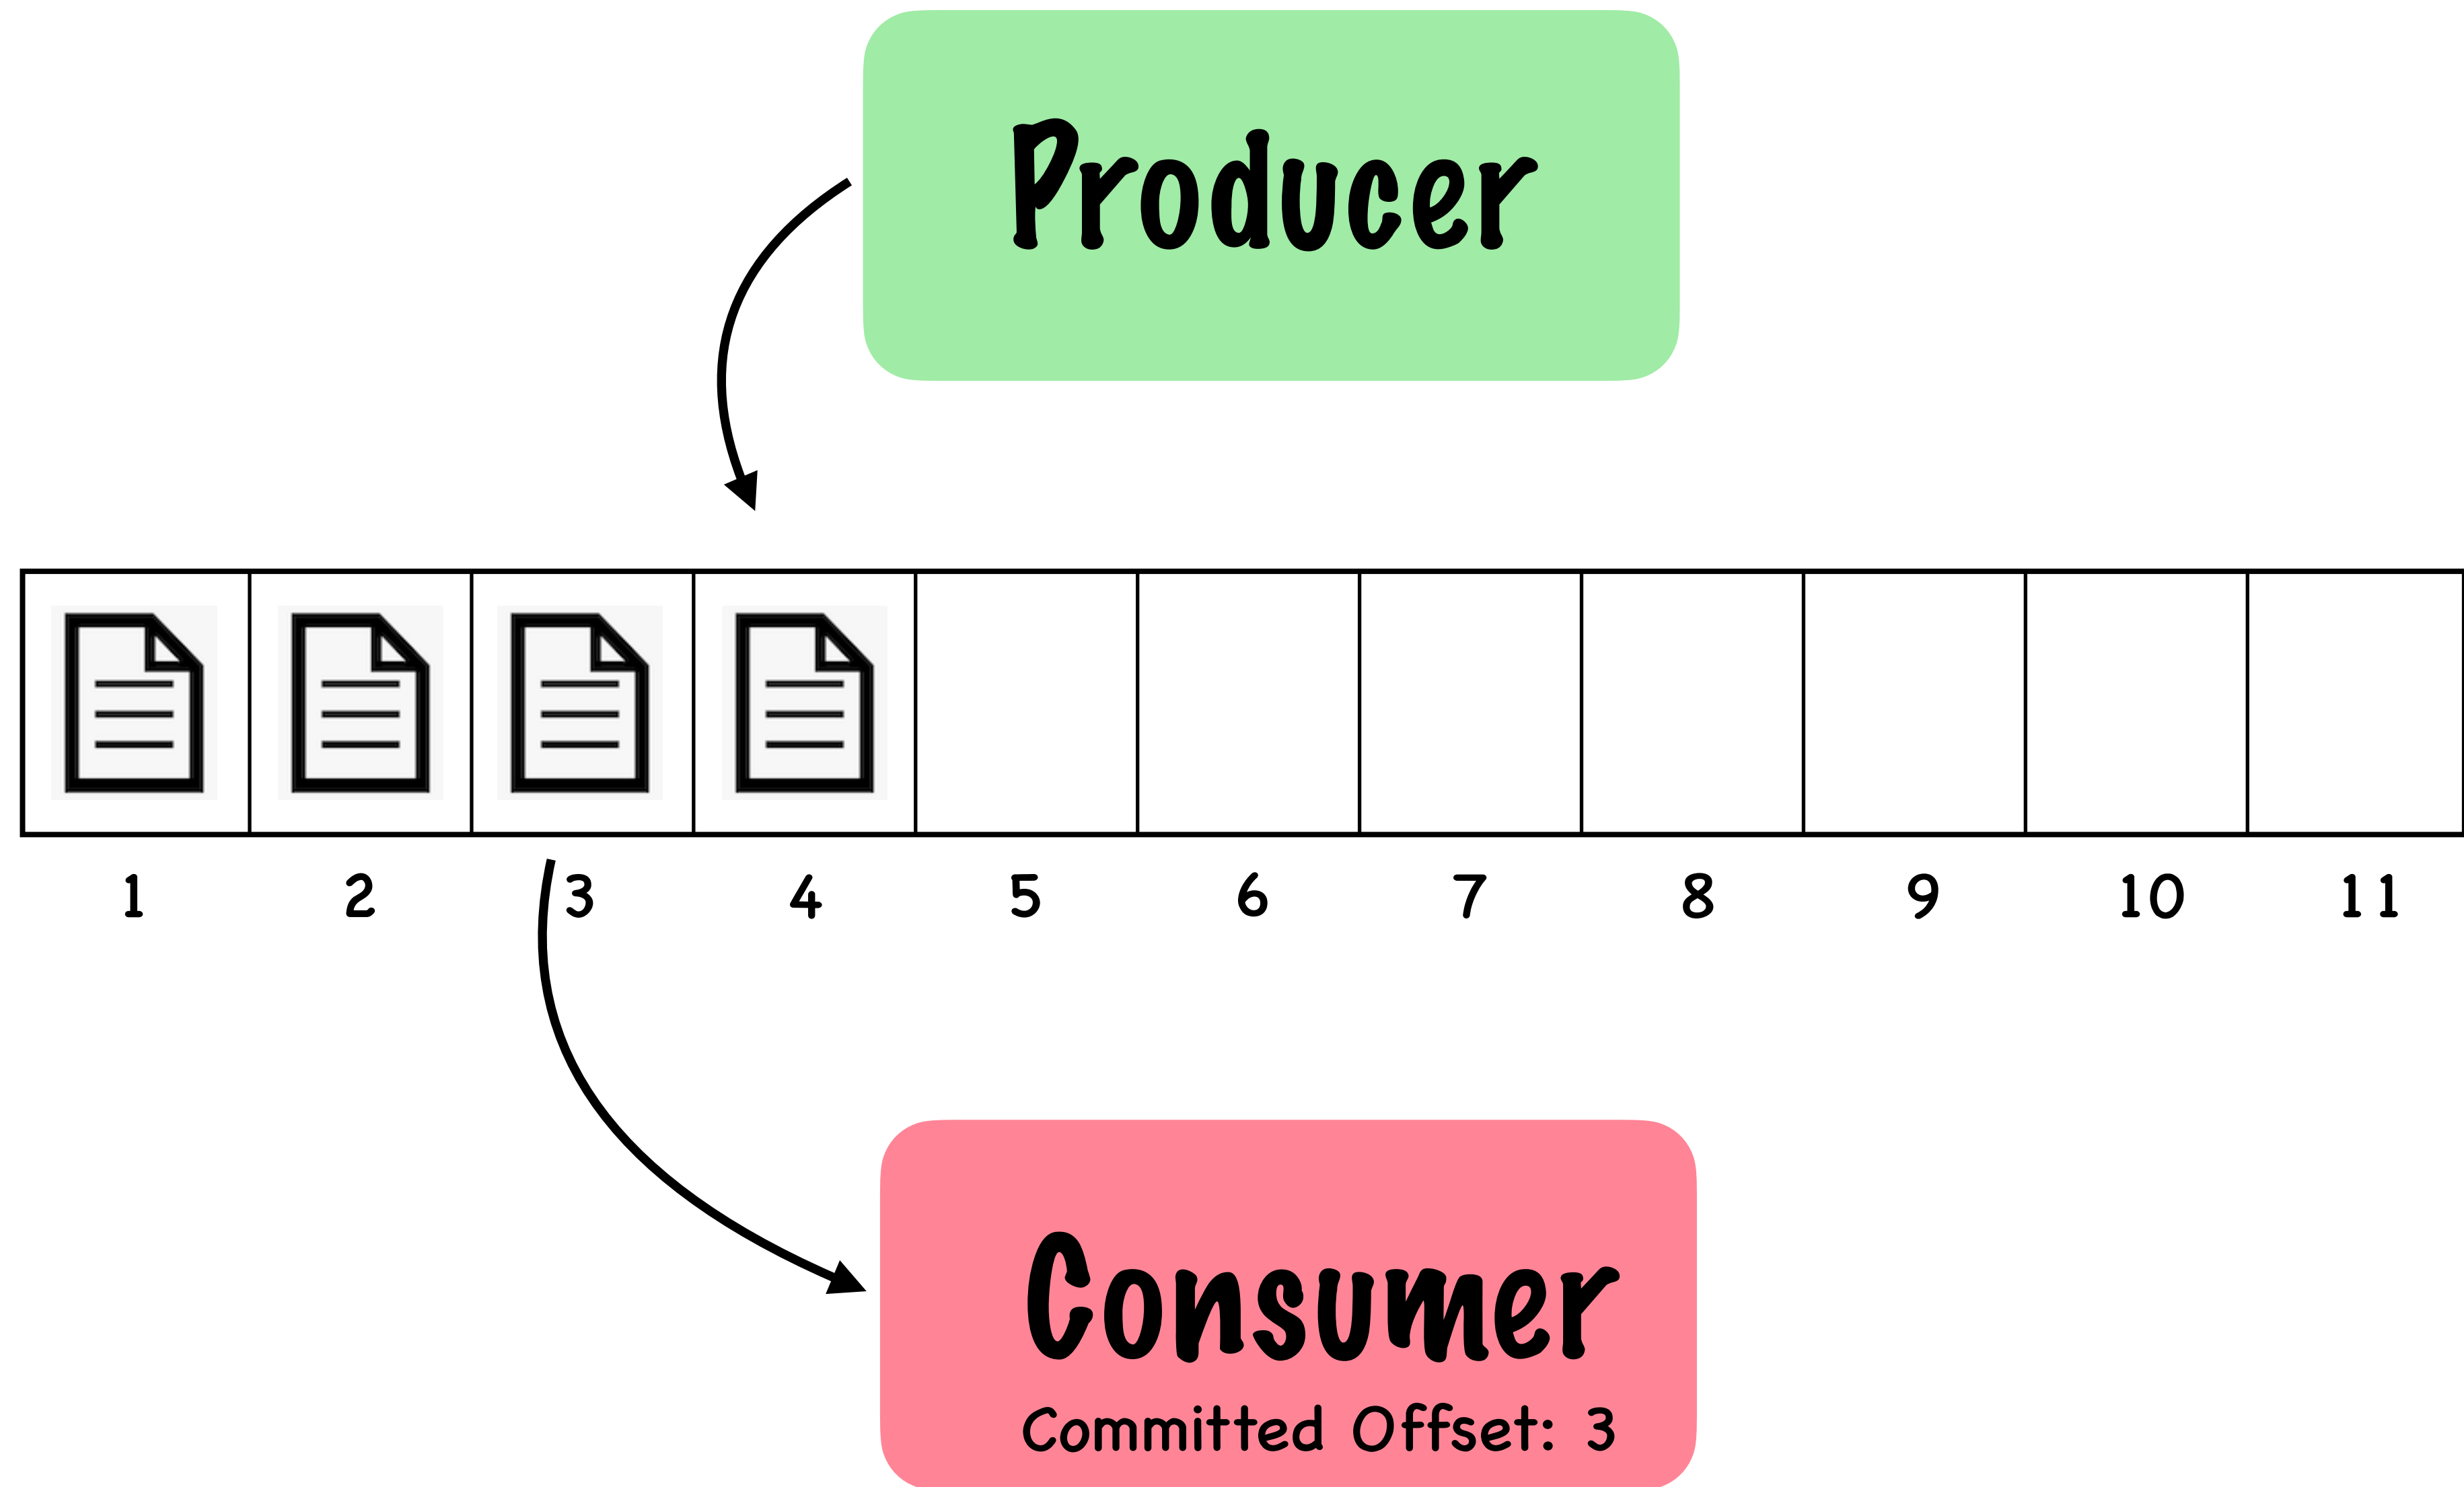















Producer



~~**Consumer**~~

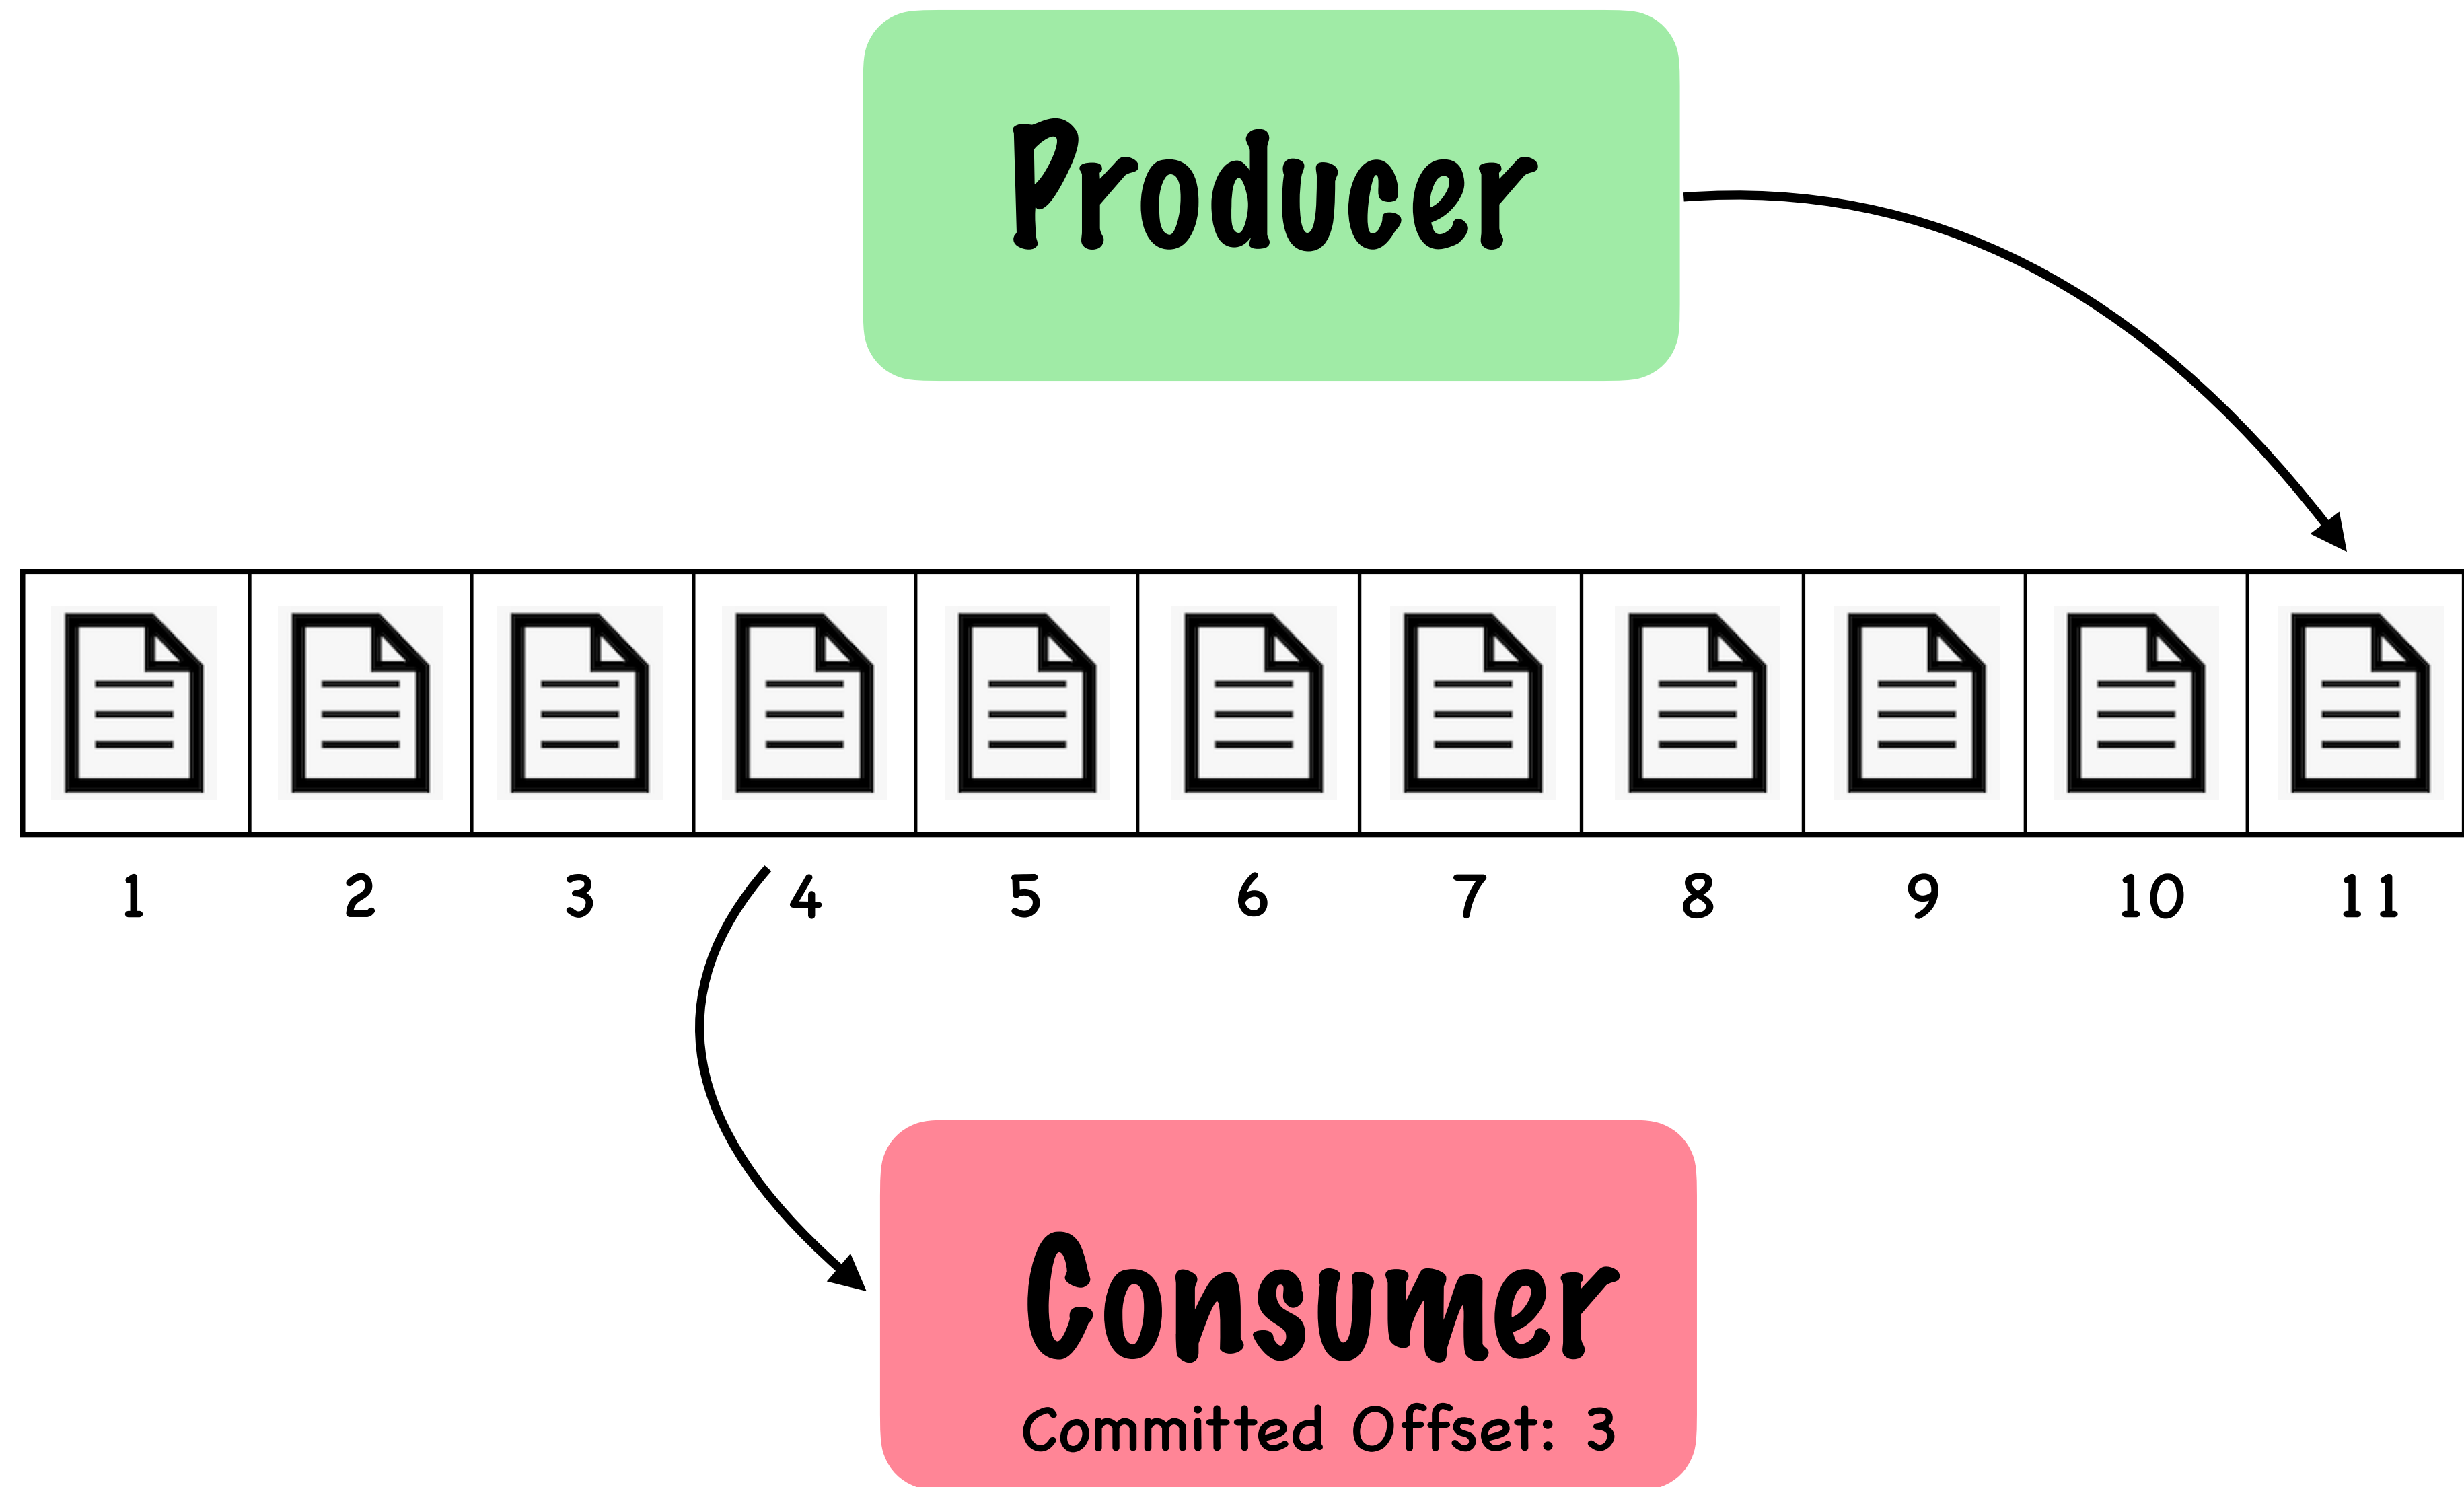
~~Committed Offset: 3~~

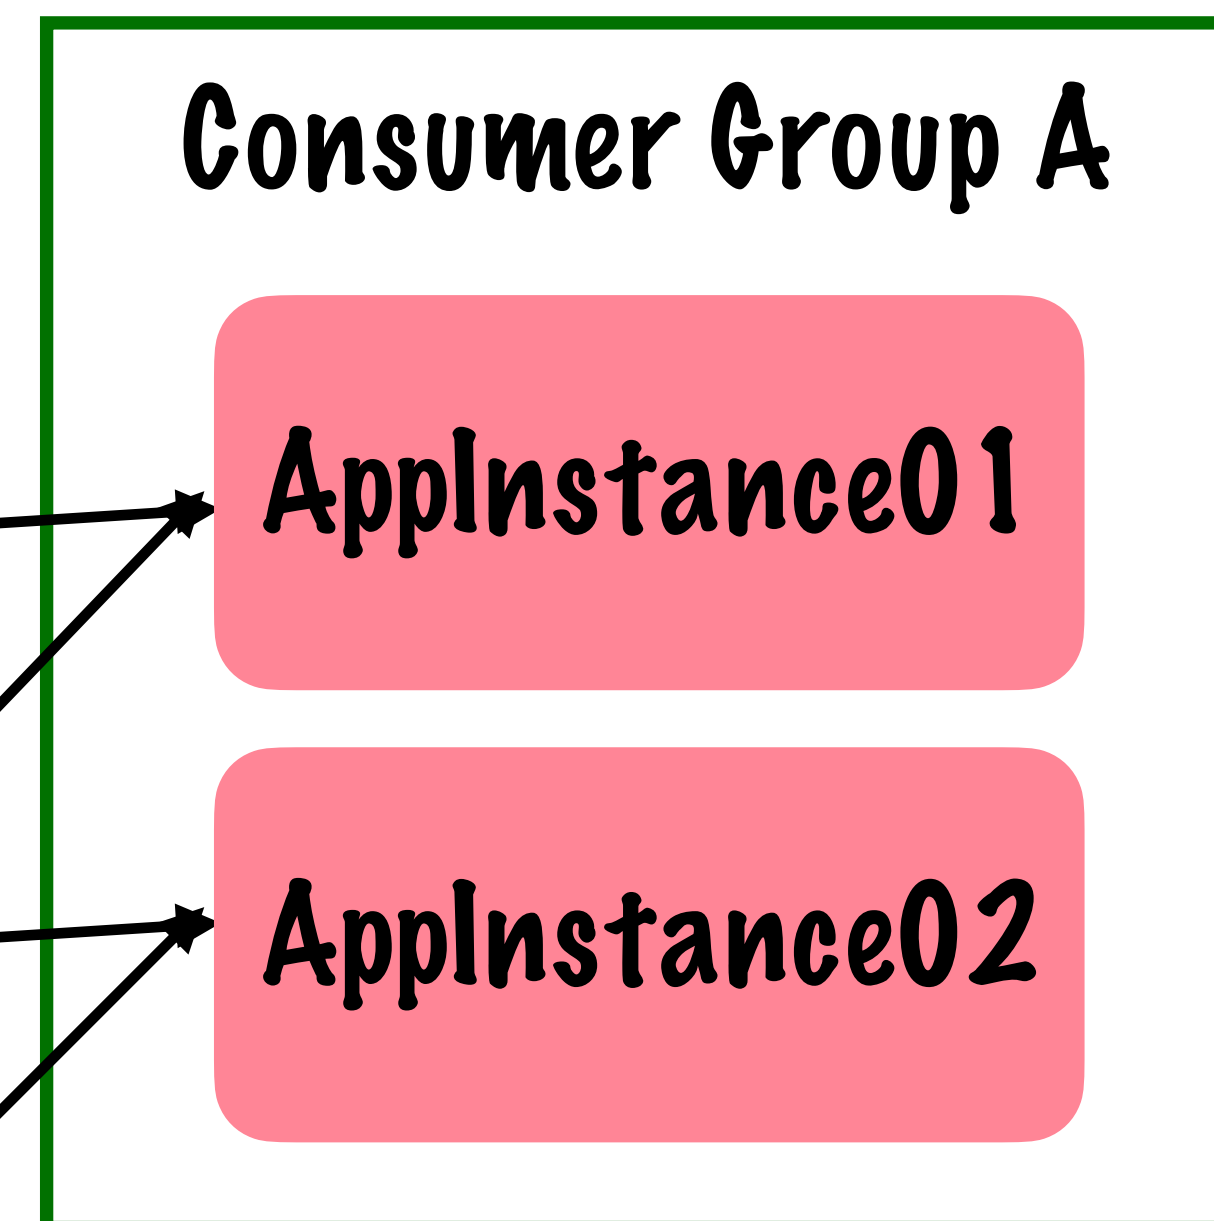

Producer




Consumer

Committed Offset: 3








Consumer Group A

ApplInstance01

ApplInstance02

ApplInstance03

ApplInstance04



Consumer Group A

ApplInstance01

ApplInstance02

Consumer Group B

OtherApp01

OtherApp02

Python Kafka Clients

confluent-kafka

<https://github.com/confluentinc/confluent-kafka-python>

kafka-python

<https://github.com/dpkp/kafka-python>

aiokafka

<https://github.com/aio-libs/aiokafka>

Producing data to a Kafka topic

```
from confluent_kafka import Producer

my_key = "a001"
my_val = "Hello World"

producer = Producer({"bootstrap.servers": "localhost:29092"})

producer.produce("my_topic", key=my_key, value=my_val)
```

Consuming data from a Kafka topic

```
from confluent_kafka import Consumer

consumer = Consumer({"bootstrap.servers": "localhost:29092",
                    "group.id": "my_consumer_group"})

consumer.subscribe(["my_topic"])

while True:
    event = consumer.poll(1.0)
    if event is None:
        pass
    else:
        val = event.value().decode("utf-8")
        print(val)
```

Stream Processing

Input

Stream Processing



Output

Stream Processing

Stateless Operations

Filter

```
producer = Producer(...)
consumer = Consumer(...)
consumer.subscribe(["input_topic"])

while True:
    event = consumer.poll(1.0)
    if check_predicate(event):
        producer.produce("output_topic", value=event)
```

Stream Processing

Stateless Operations

Branch

```
producer = Producer(...)
consumer = Consumer(...)
consumer.subscribe(["input_topic"])

while True:
    event = consumer.poll(1.0)
    if check_predicate_a(event):
        producer.produce("output_a", value=event)
    elif check_predicate_b(event):
        producer.produce("output_b", value=event)
    else:
        producer.produce("output_other", value=event)
```


Stream Processing

Stateless Operations

Map

```
producer = Producer(...)
consumer = Consumer(...)
consumer.subscribe(["input_topic"])

def my_func(event):
    # some process we want done on all events

while True:
    my_func(consumer.poll(1.0))
```

Stream Processing

Stateful Operations

Count

```
...
counts = {}
while True:
    event = consumer.poll(1.0)
    if event.key() in counts:
        counts[event.key()] += 1
    else:
        counts[event.key()] = 1
    producer.produce("output", key=event.key(),
                     value=counts[event.key()])
```

Stream Processing

Stateful Operations

Sum

```
...
sums = {}
while True:
    event = consumer.poll(1.0)
    if event.key() in sums:
        sums[event.key()] += event.value()
    else:
        sums[event.key()] = event.value()
    producer.produce("output", key=event.key(),
                     value=sums[event.key()])
```

Stream Processing

Stateful Operations

Aggregate

```
...
aggs = {}
while True:
    event = consumer.poll(1.0)
    aggs[event.key()] = agg_func(aggs[event.key()], event.value())
    producer.produce("output", key=event.key(),
                     value=aggs[event.key()])
```

Stateful Operations

Join

```
# a_dict, b_dict, c_dict, and consumers declared above
while True:
    c_event = None
    a_event = consumer_a.poll(1.0)
    a_dict[a_event.key()] = a_event
    if a_event.key() in b_dict and a_event.key() not in c_dict:
        c_event = join_func(a_event, b_dict[a_event.key()])
        c_dict[a_event.key()] = c_event

    b_event = consumer_b.poll(1.0)
    b_dict[b_event.key()] = b_event
    if b_event.key() in a_dict and b_event.key() not in c_dict:
        c_event = join_func(a_dict[b_event.key()], b_event)
        c_dict[b_event.key()] = c_event

    if c_event is not None:
        producer.produce(c_topic, key=c_event.key(), value=c_event)
```

Stateful Operations

Join

```
# a_dict, b_dict, c_dict, and consumers declared above
while True:
    c_event = None
    a_event = consumer_a.poll(1.0)
    a_dict[a_event.key()] = a_event
    if a_event.key() in b_dict and a_event.key() not in c_dict:
        c_event = join_func(a_event, b_dict[a_event.key()])
        c_dict[a_event.key()] = c_event

    b_event = consumer_b.poll(1.0)
    b_dict[b_event.key()] = b_event
    if b_event.key() in a_dict and b_event.key() not in c_dict:
        c_event = join_func(a_dict[b_event.key()], b_event)
        c_dict[b_event.key()] = c_event

    if c_event is not None:
        producer.produce(c_topic, key=c_event.key(), value=c_event)
```

Stateful Operations

Join

```
# a_dict, b_dict, c_dict, and consumers declared above
while True:
    c_event = None
    a_event = consumer_a.poll(1.0)
    a_dict[a_event.key()] = a_event
    if a_event.key() in b_dict and a_event.key() not in c_dict:
        c_event = join_func(a_event, b_dict[a_event.key()])
        c_dict[a_event.key()] = c_event

    b_event = consumer_b.poll(1.0)
    b_dict[b_event.key()] = b_event
    if b_event.key() in a_dict and b_event.key() not in c_dict:
        c_event = join_func(a_dict[b_event.key()], b_event)
        c_dict[b_event.key()] = c_event

    if c_event is not None:
        producer.produce(c_topic, key=c_event.key(), value=c_event)
```


Stateful Operations

Join

```
# a_dict, b_dict, c_dict, and consumers declared above
while True:
    c_event = None
    a_event = consumer_a.poll(1.0)
    a_dict[a_event.key()] = a_event
    if a_event.key() in b_dict and a_event.key() not in c_dict:
        c_event = join_func(a_event, b_dict[a_event.key()])
        c_dict[a_event.key()] = c_event

    b_event = consumer_b.poll(1.0)
    b_dict[b_event.key()] = b_event
    if b_event.key() in a_dict and b_event.key() not in c_dict:
        c_event = join_func(a_dict[b_event.key()], b_event)
        c_dict[b_event.key()] = c_event

    if c_event is not None:
        producer.produce(c_topic, key=c_event.key(), value=c_event)
```

Stateful Operations

Join

```
# a_dict, b_dict, c_dict, and consumers declared above
while True:
    c_event = None
    a_event = consumer_a.poll(1.0)
    a_dict[a_event.key()] = a_event
    if a_event.key() in b_dict and a_event.key() not in c_dict:
        c_event = join_func(a_event, b_dict[a_event.key()])
        c_dict[a_event.key()] = c_event

    b_event = consumer_b.poll(1.0)
    b_dict[b_event.key()] = b_event
    if b_event.key() in a_dict and b_event.key() not in c_dict:
        c_event = join_func(a_dict[b_event.key()], b_event)
        c_dict[b_event.key()] = c_event

    if c_event is not None:
        producer.produce(c_topic, key=c_event.key(), value=c_event)
```

Stateful Operations

Join

```
# a_dict, b_dict, c_dict, and consumers declared above
while True:
    c_event = None
    a_event = consumer_a.poll(1.0)
    a_dict[a_event.key()] = a_event
    if a_event.key() in b_dict and a_event.key() not in c_dict:
        c_event = join_func(a_event, b_dict[a_event.key()])
        c_dict[a_event.key()] = c_event

    b_event = consumer_b.poll(1.0)
    b_dict[b_event.key()] = b_event
    if b_event.key() in a_dict and b_event.key() not in c_dict:
        c_event = join_func(a_dict[b_event.key()], b_event)
        c_dict[b_event.key()] = c_event

    if c_event is not None:
        producer.produce(c_topic, key=c_event.key(), value=c_event)
```

Stateful Operations

Join

```
# a_dict, b_dict, c_dict, and consumers declared above
while True:
    c_event = None
    a_event = consumer_a.poll(1.0)
    a_dict[a_event.key()] = a_event
    if a_event.key() in b_dict and a_event.key() not in c_dict:
        c_event = join_func(a_event, b_dict[a_event.key()])
        c_dict[a_event.key()] = c_event

    b_event = consumer_b.poll(1.0)
    b_dict[b_event.key()] = b_event
    if b_event.key() in a_dict and b_event.key() not in c_dict:
        c_event = join_func(a_dict[b_event.key()], b_event)
        c_dict[b_event.key()] = c_event

    if c_event is not None:
        producer.produce(c_topic, key=c_event.key(), value=c_event)
```

Stateful Operations

Join

```
# a_dict, b_dict, c_dict, and consumers declared above
while True:
    c_event = None
    a_event = consumer_a.poll(1.0)
    a_dict[a_event.key()] = a_event
    if a_event.key() in b_dict and a_event.key() not in c_dict:
        c_event = join_func(a_event, b_dict[a_event.key()])
        c_dict[a_event.key()] = c_event

    b_event = consumer_b.poll(1.0)
    b_dict[b_event.key()] = b_event
    if b_event.key() in a_dict and b_event.key() not in c_dict:
        c_event = join_func(a_dict[b_event.key()], b_event)
        c_dict[b_event.key()] = c_event

    if c_event is not None:
        producer.produce(c_topic, key=c_event.key(), value=c_event)
```

Stateful Operations

Join

```
# a_dict, b_dict, c_dict, and consumers declared above
while True:
    c_event = None
    a_event = consumer_a.poll(1.0)
    a_dict[a_event.key()] = a_event
    if a_event.key() in b_dict and a_event.key() not in c_dict:
        c_event = join_func(a_event, b_dict[a_event.key()])
        c_dict[a_event.key()] = c_event

    b_event = consumer_b.poll(1.0)
    b_dict[b_event.key()] = b_event
    if b_event.key() in a_dict and b_event.key() not in c_dict:
        c_event = join_func(a_dict[b_event.key()], b_event)
        c_dict[b_event.key()] = c_event

    if c_event is not None:
        producer.produce(c_topic, key=c_event.key(), value=c_event)
```

Stream Processing Challenges


State

Scale

Python Event Streaming Libraries



Python Event Streaming Libraries



<https://github.com/faust-streaming/faust>



<https://github.com/bytewax/bytewax>



<https://github.com/quixio/quix-streams>